

Parallel Computing with GPUs: Visual Studio Guide for CUDA

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



The
University
Of
Sheffield.

Compiling a CUDA program

- ❑ CUDA C Code is compiled using **nvcc** e.g.
- ❑ Will compile host AND device code to produce an executable

```
nvcc -o example example.cu
```

- ❑ We will be using Visual Studio to build our CUDA code so we will not need to compile at the command line (unless you are running on ShARC)

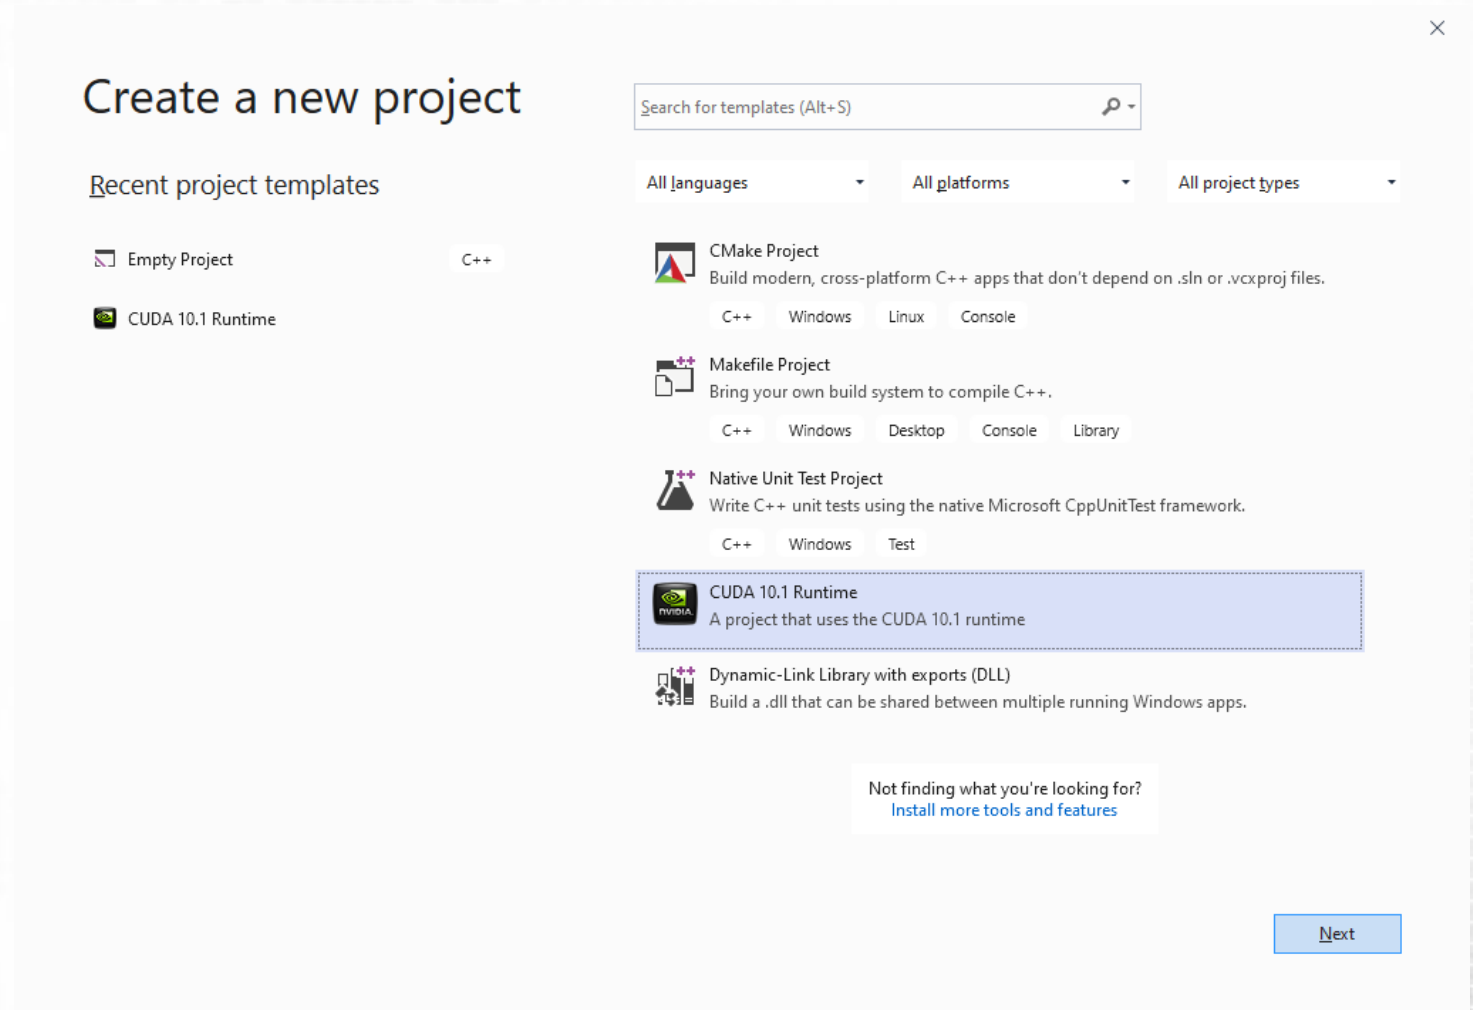
Creating a CUDA Project

❑ Create New CUDA Project

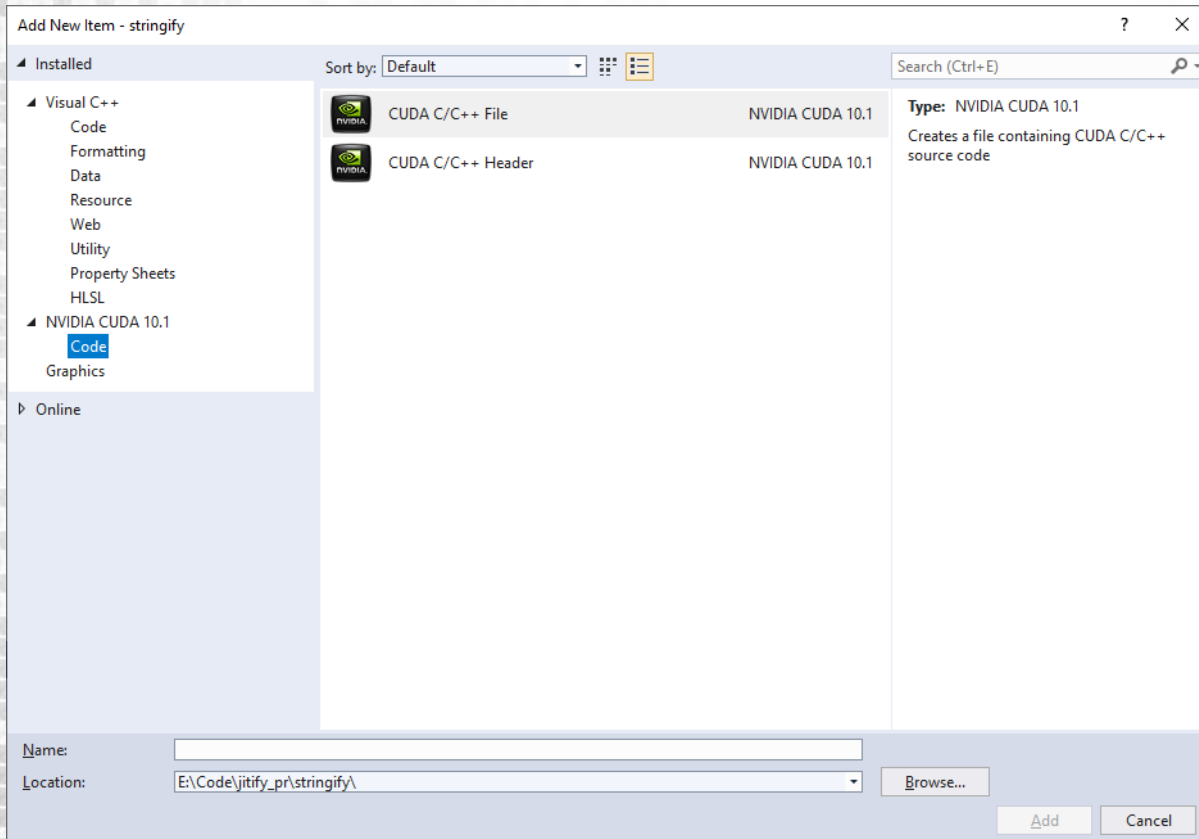
❑ Select *NVIDIA -> CUDA 11.1*

❑ This will create a project with a default kernels.cu file containing a basic vector addition example

Preferred Method!



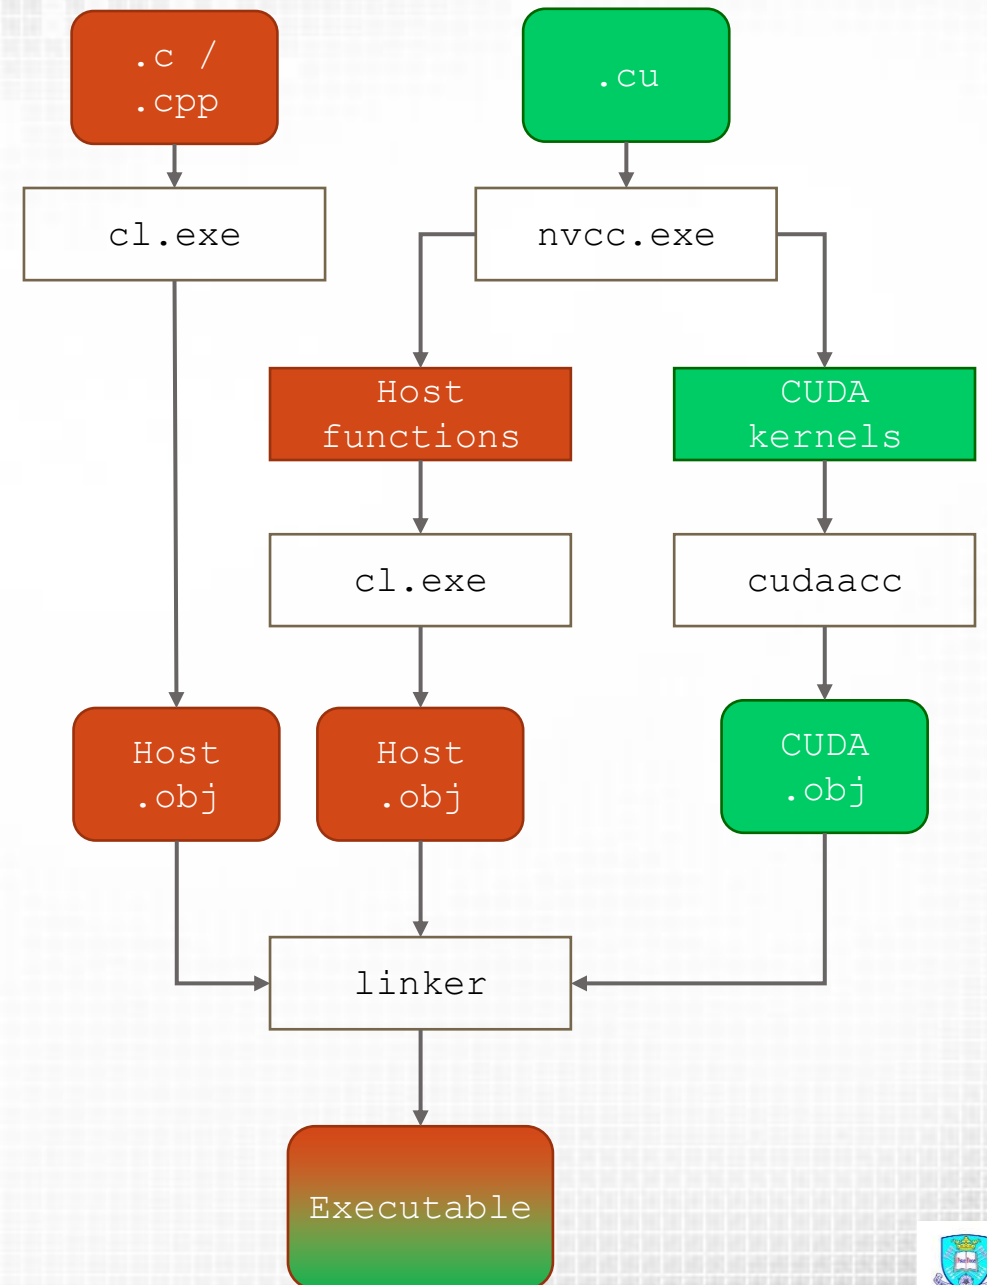
Adding a CUDA source file



- ❑ Alternatively add a CUDA source file to an existing application
- ❑ If you do this you must modify the project properties to include CUDA build customisations
 - ❑ http://developer.download.nvidia.com/compute/cuda/6_5/rel/docs/CUDA_Getting_Started_Windows.pdf (section 3.4)

Compilation

- ❑ CUDA source file (* .cu) are compiled by `nvcc`
- ❑ An existing `cuda.rules` file creates property page for CUDA source files
 - ❑ Configures `nvcc` in the same way as configuring the C compiler
 - ❑ Options such as optimisation and include directories can be inherited from project defaults
- ❑ C and C++ files are compiled with `cl` (MSVCC compiler)



Device Versions

- ❑ Different generations of NVIDIA hardware have different compatibility
 - ❑ In the last lecture we saw product families and chip variants
 - ❑ These are classified by **CUDA compute versions**
- ❑ Compilation normally builds for CUDA compute version 35
 - ❑ See Project Properties, *CUDA C/C++ Device->Code Generation*
 - ❑ Default value is “`compute_35, sm_35`”
 - ❑ Any hardware with greater than the compiled compute version can execute the code (backwards compatibility)
- ❑ You can build for multiple versions using separator
 - ❑ E.g. “`compute_35, sm_35; compute_60, sm_60`”
 - ❑ This will increase build time and execution file size
 - ❑ Runtime will select the best version for your hardware

Device Versions of Available GPUs

- ❑ Diamond Machines

- ❑ Pascal Architecture

- ❑ `compute_60,sm_60;`

CUDA Properties

main.cu Property Pages

Configuration: Debug Platform: x64 Configuration Manager...

- Configuration Properties
 - General
 - CUDA C/C++
 - Common
 - Device**
 - Host
 - Command Line

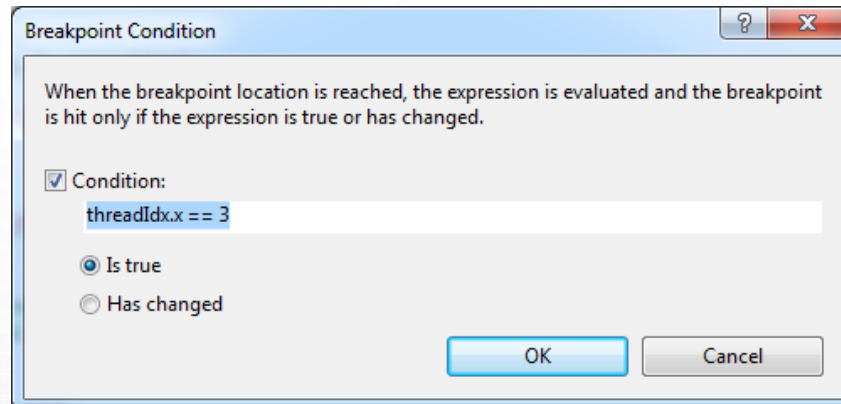
Interleave source in PTX	No
Code Generation	compute_75,sm_75;compute_75,compute_75
Generate GPU Debug Information	Yes (-G)
Generate Line Number Information	No
Max Used Register	0
Verbose PTXAS Output	No

Interleave source in PTX
Insert source code in generated PTX.

OK Cancel Apply

Debugging

- ❑ NSIGHT is a GPU debugger for debugging GPU kernel code
 - ❑ It does not debug breakpoints in host code
- ❑ To launch select insert a breakpoint and select NSIGHT-> Start CUDA Debugging
 - ❑ You must be in the debug build configuration.
 - ❑ When stepping all warps except the debugger focus will be paused
- ❑ Use conditional breakpoints to focus on specific threads
 - ❑ Right click on break point and select Condition



Error Checking

- ❑ `cudaError_t`: enumerator for runtime errors
 - ❑ Can be converted to an error string (`const char *`) using `cudaGetErrorString(cudaError_t)`
- ❑ Many host functions (e.g. `cudaMalloc`, `cudaMemcpy`) return a `cudaError_t` which can be used to handle errors gracefully

```
cudaError_t cudaStatus;  
  
cudaStatus = cudaMemcpy(dev_a, a, size * sizeof(int), cudaMemcpyHostToDevice);  
if (cudaStatus != cudaSuccess) {  
    //handle error  
}
```

- ❑ Kernels do not return an error but if one is raised it can be queried using the `cudaGetLastError()` function

```
addKernel<<<1, size>>>(dev_c, dev_a, dev_b);  
cudaStatus = cudaGetLastError();
```